

# Parallel Programming using MPI

Lecture 3  
January 12, 2026

# Parallel Programming Models

- Shared memory
- Distributed memory

# Parallel Programming Models (Recap)

## Shared memory programming – OpenMP

- Shared address space
- Implicit communication



Legend:  
Blue square: Cache  
Yellow square: Core  
Green square: Thread


## Distributed memory programming – MPI

- Distinct address space
- Explicit communication



Legend:  
Blue square: Cache  
Yellow square: Core  
Green square: Process

# Process (Recap)




## Process:

- \* An instance of a program execution.
- \* The execution context of a running program... i.e. the resources associated with a program's execution.



# Memory Layout?



# Simple MPI Code

```
#include <stdio.h>
#include "mpi.h"

int main(int argc, char *argv[])
{
    // initialize MPI
    MPI_Init (&argc, &argv);

    printf ("Hello, world!\n");

    // done with MPI
    MPI_Finalize();
}

~
```

# MPI Code Execution Steps

- Compile
  - `mpicc -o program.x program.c`
- Execute
  - `mpirun -np 1 ./program.x` (`mpiexec -np 1 ./program.x`)
    - Runs 1 process on the launch/login node
  - `mpirun -np 6 ./program.x`
    - Runs 6 processes on the launch/login node

# Output – Hello World

```
mpirun -np 20 ./program.x
```

```
Hello, world!
Hello, world!Hello, world!
Hello, world!
```

# Beowulf PC Clusters: Breaking the Cost Barrier to High End Application Computing

Thomas Sterling, PhD

"Leader of NASA BEOWULF Project"

Senior Staff Scientist; HPC Systems Group, Jet Propulsion Lab  
Visiting Professor, California Institute of Technology

This exciting presentation emphasizes the economical opportunity of networking low cost personal computers (PCs) to solve science and engineering problems. Such PC networks permit parallel or distributed computing in an effective and efficient manner. Various high performance computing applications using personal computer networks are addressed in terms of science performed and performance achieved. Performance metrics discussed include speedup vs. costs vs. scalability. Such experience has evolved with Dr. Sterling's NASA BEOWULF project and associated BEOWULF-class computing. Come and discuss your applications using a PC network for high performance computation!

# **FROM TOYS TO TERAFLOPS: BRIDGING THE BEOWULF GAP**

**Thomas Sterling  
Daniel F. Savarese**

CENTER FOR ADVANCED COMPUTING RESEARCH,  
CALIFORNIA INSTITUTE OF TECHNOLOGY, PASADENA,  
CALIFORNIA, U.S.A.

## **Summary**

Do-it-yourself supercomputing has emerged as a solution to cost-effectively sustain the computational demands of the scientific research community. Despite some of the successes of this approach, represented by Beowulf-class computing, it has limitations that need to be recognized as well as problems that need to be resolved in order to extend its scope of applicability. While the performance of hardware incorporated into these systems has continued to improve at a remarkable rate, enabling the execution of steadily larger and more compute-intensive applications, the software environment of the machines has seen little to no improvement or evolution. The authors find

## **1 A Crossroads of Computing**

Over the past 5 years, the do-it-yourself supercomputing phenomenon has grown in popularity as the ranks of high performance computer vendors have winnowed and the performance of commodity microprocessors has steadily improved. With the maturation of no-cost operating systems and compilers into high performance commercial-quality software, in addition to the establishment of standard parallel programming models, the case has become more compelling than ever for research organizations to assemble their own parallel computing platforms. The high cost of commercial supercomputers and the oversubscription of resources at supercomputing centers have forced many researchers to resort to their own means in order to run computational simulations. Ensembles of networked desktop computers have become a liberating force for computational science, enabling supercomputing to be performed in the figurative backyard.

We have entered a crossroads, where the notion of supercomputing is losing, or at least redefining, its meaning, as it becomes subsumed by the more general notion of parallel computing. The prefix “super” is applied to a noun to highlight its extraordinary nature, elevating it above the norm. Commodity-clustered computing is anything but extraordinary. Yet, there is little to distinguish

# CSEWS\*





```
Architecture:          x86_64  
CPU op-mode(s):       32-bit, 64-bit  
Byte Order:           Little Endian  
CPU(s):               12  
On-line CPU(s) list: 0-11  
Thread(s) per core:   2  
Core(s) per socket:   6  
Socket(s):            1  
NUMA node(s):         1  
Vendor ID:            GenuineIntel  
CPU family:           6  
Model:                158  
Model name:           Intel(R) Core(TM) i7-8700 CPU @ 3.20GHz  
Stepping:              10  
CPU MHz:              900.353  
CPU max MHz:          4600.0000  
CPU min MHz:          800.0000  
BogoMIPS:              6384.00  
Virtualization:        VT-x  
L1d cache:             32K  
L1i cache:             32K  
L2 cache:              256K  
L3 cache:              12288K  
NUMA node0 CPU(s):     0-11
```

*lscpu*

processor	:	0
processor	:	1
processor	:	2
processor	:	3
processor	:	4
processor	:	5
processor	:	6
processor	:	7
processor	:	8
processor	:	9
processor	:	10
processor	:	11

# Process Placement


`mpirun -np 6 ./program.x`




# Unmanaged Beowulf Cluster

`mpirun -np 6 ./program.x`

Running  
program  
(others'  
jobs)



Your  
application



# Execution Parameters





MPI process  
MPI process

Number of nodes = 4

Processes per node (ppn) = 2



# Process Placement on Unmanaged Cluster

`mpirun -np 8 -f hostfile ./program.x`




nodes=4

ppn=2



# Multiple Processes Per Node



Run examples/cpi (mpirun -np 8 <path to examples/cpi>)

# MPI Code Execution Steps

- Compile
  - `mpicc -o program.x program.c`
- Execute
  - `mpirun -np 1 ./program.x`
    - Runs 1 process on the launch/login node
  - `mpirun -np 6 ./program.x`
    - Runs 6 processes on the launch/login node
  - `mpirun -np 6 -f hostfile ./program.x`
    - Runs 6 processes on the nodes specified in the hostfile

# Hostfile – Options

172.27.19.1

172.27.19.2

172.27.19.3

172.27.19.4

172.27.19.1:1

172.27.19.2:1

172.27.19.3:1

172.27.19.4:1

172.27.19.1:4


172.27.19.2:4

172.27.19.3:4

172.27.19.4:4

# Round-robin Placement in Hydra


172.27.19.1  
172.27.19.2  
172.27.19.3  
172.27.19.4



mpiexec –np 8 –f hostfile ./program.x

# Sequential Placement in Hydra


172.27.19.1:2  
172.27.19.2:2  
172.27.19.3:2  
172.27.19.4:2



mpiexec –np 8 –f hostfile ./program.x


# Sequential Placement in Hydra

172.27.19.1:2  
172.27.19.2:2  
172.27.19.3:1  
172.27.19.4:3



mpiexec –np 8 –f hostfile ./program.x

# Parallel Sum Execution




parallelsum parallelsum parallelsum parallelsum  
.x .x .x .x

**mpirun –np 4 –f hostfile ./parallelsum.x**

Communication  
rounds (log P)


# Parallel Sum Execution (8 processes)



Q: Is the depicted **communication pattern** correct for optimized algorithm of parallel sum?

A: No


# Parallel Sum Execution (8 processes)



Communication  
rounds ( $\log P$ )


Round 1

Partial sums at processes 0, 2, 4, 6



Round 2

# Parallel Sum Execution




Partial sums at processes 0, 4



Final sum at process 0


# Parallel Sum Execution (Sequential placement)



Partial sums at processes 0, 2, 4, 6



# Parallel Sum Execution



Communication  
rounds ( $\log P$ )

Partial sums at processes 0, 4



Round 3  
(2 hops)

Final sum at process 0

# Effect of Process Placement (Takeaway from the previous examples)

- Communication pattern is same
  - because the parallel algorithm is same
- The communication time may not be same
  - depends on the placement of communicating processes
  - also depends on the other concurrent communications

# MPI Code Execution Options


Same host

- `mpirun -np 6 ./program.x`

Multiple hosts


- `mpirun -np 6 -f hostfile ./program.x`
  - Round robin process placement
    - `mpirun -np 6 --hosts csews1,csews2 ./program.x`
  - Sequential/contiguous placement
    - `mpirun -np 6 --hosts csews1:3,csews2:3 ./program.x`

# Access Latency (Intel KNL)




Simulating the Network Activity of  
Modern Manycores, IEEE Access 2019

# Cores




Simulating the Network Activity of  
Modern Manycores, IEEE Access 2019

# Parallel Execution




# Parallel Execution




mpirun –np 8 –f hostfile ./parallel.x


# Parallel Sum (Optimized - Recap)



(a)



(b)




(c)



(d)

Source: GGKK Chapter 5

# Sum of Numbers Speedup (Optimized)



Parallel algorithm

- Compute in parallel
- For log P steps/rounds (multiple **concurrent** communications within a step/round)
  - Send partial result to a neighbour
  - Compute partial sum

$$S_P = \frac{N}{\frac{N}{P} + 2\log P}$$

$$E_P = \frac{1}{\frac{2P\log P}{N} + 1}$$

# Homework

Analyze the communication endpoints for the optimized algorithm of parallel sum for a 2D mesh network topology for ppn=1, ppn=2 (total 8 processes on 4 nodes)

# Getting Started

```
#include <stdio.h>
#include "mpi.h"

int main(int argc, char *argv[])
{
    // initialize MPI
    MPI_Init (&argc, &argv);

    printf ("Hello, world!\n");

    // done with MPI
    MPI_Finalize();
}
~
~
```

- gather information about the parallel job
- set up internal library state
- prepare for communication

Initialization

Finalization

# MPI Process Identification

Rank of a process

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {

    // Initialize the MPI environment
    MPI_Init(NULL, NULL);

    // Get the number of processes
    int size;
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    // Get the rank of the process
    int rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    // Print off a hello world message
    printf("Hello I am rank %d out of %d processes\n", rank, size);

    // Finalize the MPI environment.
    MPI_Finalize();
}
```

Global  
communicator

Total number  
of processes

# Entities

**Communicator** (communication handle)

- Defines the scope
- Specifies communication context

Process


- Belongs to a group
- Identified by a rank within a group

Identification

- MPI\_Comm\_size – total number of processes in communicator
- MPI\_Comm\_rank – rank in the communicator

# MPI\_COMM\_WORLD

- Required in every MPI communication
- Process identified by rank/id



# Communicator

- Communication handle among a group/collection of processes
- Representative of communication domain
- Associated with a context ID (in MPICH)
- Predefined:
  - MPI\_COMM\_WORLD
  - MPI\_COMM\_SELF

# Output

```
// get number of tasks  
MPI_Comm_size (MPI_COMM_WORLD, &numtasks);  
  
// get my rank  
MPI_Comm_rank (MPI_COMM_WORLD, &rank);  
  
printf ("Hello I'm rank %d of %d processes\n", rank, numtasks);
```


E1: mpiexec -n 4 -hosts host1,host2,host3,host4 ./exe

E2: mpiexec -n 4 -hosts host1:2,host2:2,host3:2,host4:2 ./exe

E3: mpiexec -n 8 -hosts host1,host2,host3,host4 ./exe

E4: mpiexec -n 8 -hosts host1:1,host2:2,host3:3,host4:4 ./exe

# Host



A screenshot of a terminal window titled "2. pmalakar@172.27.19.10". The window contains the following MPI C code:

```
#include <stdio.h>
#include "mpi.h"

int main(int argc, char *argv[])
{
    int numtasks, rank, len;
    char hostname[MPI_MAX_PROCESSOR_NAME];

    // initialize MPI
    MPI_Init (&argc, &argv);

    // get number of tasks
    MPI_Comm_size (MPI_COMM_WORLD, &numtasks);

    // get my rank
    MPI_Comm_rank (MPI_COMM_WORLD, &rank);


    // this one is obvious
    MPI_Get_processor_name (hostname, &len);

    printf ("Number of tasks=%d My rank=%d Running on %s\n", numtasks, rank, hostname);

    // done with MPI
    MPI_Finalize();
}
```

# Process Launch

Memory




Reading (optional):

[A Scalable Process-Management Environment for Parallel Programs](#)

# Parallel Environment


- Job scheduler
  - Resources to run the parallel job
- Process manager
  - Starts and stops the processes
- Parallel library
  - Application calls for communication

# Multiprocess Daemon



Butler et al., EuroPVM/MPI2000

# Process Management Setup



Parallel program library (e.g. MPI)

Process management interface (PMI)

Resource manager/  
Job scheduler/  
Process Manager

Reading (optional):

[PMI: A Scalable Parallel Process-Management Interface for Extreme-Scale Systems](#)


# Process Manager

- Start and stop processes in a **scalable** way
- Setup communication channels for parallel processes
- stdin/stdout/stderr forwarding
  - ./a.out 2>efile

# Process Management Interface


- Interface between process manager and MPI library
- Processes can exchange information about peers by querying PMI
- Uses key-value store for process-related data

# Process Launch



# Process Launch

■ Memory




# Hydra Process Manager

- A process management system for starting parallel jobs
- Uses existing daemons (viz. ssh) to start MPI processes
- Automatically detects resource managers if any and interacts with them
- `$ mpiexec ./app`
  - Hydra gets information about allocated resources and launches processes
- Passes environment variables from the shell on which `mpiexec` is launched to the launched processes

There are others – gforker, slurm, etc.

# mpiexec

Memory



```
mpiexec -n 4 --hosts host1,host2,host3,host4 ./exe
```

# Launch Node

`mpiexec -np 8 -hosts host1:3,host2:3,host3:3 ./exe`


```
pmalakar 17952 17943 0 09:41 ?      00:00:00 /usr/lib/openssh/sftp-server
pmalakar 20853 16203 0 10:20 pts/1    00:00:00 mpiexec -np 8 -hosts 172.27.19.2 3 172.27.19.3 3 172.27
.19.4 3 ./IMB-MPI1 AllReduce
pmalakar 20854 20853 0 10:20 ?      00:00:00 /users/faculty/pmalakar/mpich-3.2.1-install/bin/hydra_p
mi_proxy --control-port 172.27.19.2:46385 --rmk user --launcher ssh --demux poll --pgid 0 --retries 10
--usize -2 --proxy-id 0
pmalakar 20855 20853 0 10:20 ?      00:00:00 /usr/bin/ssh -x 172.27.19.3 "/users/faculty/pmalakar/mp
ich-3.2.1-install/bin/hydra_pmi_proxy" --control-port 172.27.19.2:46385 --rmk user --launcher ssh --dem
ux poll --pgid 0 --retries 10 --usize -2 --proxy-id 1
pmalakar 20856 20853 0 10:20 ?      00:00:00 /usr/bin/ssh -x 172.27.19.4 "/users/faculty/pmalakar/mp
ich-3.2.1-install/bin/hydra_pmi_proxy" --control-port 172.27.19.2:46385 --rmk user --launcher ssh --dem
ux poll --pgid 0 --retries 10 --usize -2 --proxy-id 2
pmalakar 20857 20854 76 10:20 ?     00:00:03 ./IMB-MPI1 AllReduce
pmalakar 20858 20854 76 10:20 ?     00:00:03 ./IMB-MPI1 AllReduce
pmalakar 20859 20854 76 10:20 ?     00:00:03 ./IMB-MPI1 AllReduce
pmalakar 20861 17877 0 10:20 pts/4   00:00:00 ps -aef
```

# Compute Node Processes

```
pmalakar 8756 8728 0 10:18 pts/0    00:00:00 -bash
pmalakar 8759 8755 0 10:18 ?        00:00:00 /usr/lib/openssh/sftp-server
root     8781 1123 0 10:20 ?        00:00:00 sshd: pmalakar [priv]
pmalakar 8845 8781 0 10:20 ?        00:00:00 sshd: pmalakar@notty
pmalakar 8846 8845 0 10:20 ?        00:00:00 /users/faculty/pmalakar/mpich-3.2.1-install/bin/hydra_pmi_prox
y --control-port 172.27.19.2:46385 --rmk user --launcher ssh --demux poll --pgid 0 --retries 10 --usize -2 --p
roxy-id 1
pmalakar 8847 8846 99 10:20 ?      00:00:12 ./IMB-MPI1 AllReduce
pmalakar 8848 8846 99 10:20 ?      00:00:12 ./IMB-MPI1 AllReduce
pmalakar 8849 8846 99 10:20 ?      00:00:12 ./IMB-MPI1 AllReduce
```

```
pmalakar 8838 8774 0 10:20 pts/1    00:00:00 -bash
pmalakar 8841 8837 0 10:20 ?        00:00:00 /usr/lib/openssh/sftp-server
root     8851 1250 0 10:20 ?        00:00:00 sshd: pmalakar [priv]
pmalakar 8915 8851 0 10:20 ?        00:00:00 sshd: pmalakar@notty
pmalakar 8916 8915 0 10:20 ?        00:00:00 /users/faculty/pmalakar/mpich-3.2.1-install/bin/hydra_p
mi_proxy --control-port 172.27.19.2:46385 --rmk user --launcher ssh --demux poll --pgid 0 --retries 10
--usize -2 --proxy-id 2
pmalakar 8917 8916 99 10:20 ?      00:00:14 ./IMB-MPI1 AllReduce
pmalakar 8918 8916 99 10:20 ?      00:00:14 ./IMB-MPI1 AllReduce
```

# Hydra and mpiexec



Source: [wiki.mpich.org](http://wiki.mpich.org)

# MPI Process Management

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {

    // Initialize the MPI environment
    MPI_Init(NULL, NULL);

    // Get the number of processes
    int size;
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    // Get the rank of the process
    int rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    // Get the name of the processor
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;
    MPI_Get_processor_name(processor_name, &name_len);

    // Print off a hello world message
    printf("Hello I am rank %d out of %d processes\n", rank, size);

    // Finalize the MPI environment.
    MPI_Finalize();
}

~
```

Initializes  
and queries  
PMI

# csews\*

```
Architecture:          x86_64
CPU op-mode(s):       32-bit, 64-bit
Byte Order:           Little Endian
CPU(s):               12
On-line CPU(s) list: 0-11
Thread(s) per core:  2
Core(s) per socket:  6
Socket(s):            1
NUMA node(s):         1
Vendor ID:            GenuineIntel
CPU family:           6
Model:                158
Model name:           Intel(R) Core(TM) i7-8700 CPU @ 3.20GHz
Stepping:              10
CPU MHz:              900.353
CPU max MHz:          4600.0000
CPU min MHz:          800.0000
BogoMIPS:             6384.00
Virtualization:       VT-x
L1d cache:            32K
L1i cache:            32K
L2 cache:              256K
L3 cache:              12288K
NUMA node0 CPU(s):    0-11
```

*lscpu*

processor	:	0
processor	:	1
processor	:	2
processor	:	3
processor	:	4
processor	:	5
processor	:	6
processor	:	7
processor	:	8
processor	:	9
processor	:	10
processor	:	11

# Multiple Tasks Core ID

```
#define _GNU_SOURCE
#include <stdio.h>
#include <sched.h>
#include "mpi.h"

int main(int argc, char *argv[])
{
    int numtasks, rank, len, coreID;
    char hostname[MPI_MAX_PROCESSOR_NAME];

    // initialize MPI
    MPI_Init (&argc, &argv);

    // get number of tasks
    MPI_Comm_size (MPI_COMM_WORLD, &numtasks);

    // get my rank
    MPI_Comm_rank (MPI_COMM_WORLD, &rank);

    // get hostname
    MPI_Get_processor_name (hostname, &len);


    // core ID
    coreID = sched_getcpu();

    printf ("%d of %d Running on %s:%d\n", rank, numtasks, hostname, coreID);

    // done with MPI
    MPI_Finalize();
}
```

Get the  
core ID

# Process Placement



A screenshot of a terminal window titled "pmalakar@172.27.19.10". The window has a dark theme with light-colored text. It shows the command "cs633 mpirun -np 12 ./multipletaskscoreID" followed by its execution results. The results indicate 12 processes are running on nodes csews10:0 through csews10:11. The terminal also shows the user's name "pmalakar" and IP address "172.27.19.10" at the top.

```
pmalakar@172.27.19.10
Terminal Sessions View X server Tools Games Settings Macros Help
2. pmalakar@172.27.19.10 × 3. pmalakar@172.27.19.10 × +
```

```
cs633 mpirun -np 12 ./multipletaskscoreID
0 of 12 Running on csews10:0
1 of 12 Running on csews10:3
2 of 12 Running on csews10:6
6 of 12 Running on csews10:9
7 of 12 Running on csews10:1
8 of 12 Running on csews10:8
11 of 12 Running on csews10:7
3 of 12 Running on csews10:5
4 of 12 Running on csews10:11
5 of 12 Running on csews10:4
9 of 12 Running on csews10:10
10 of 12 Running on csews10:2
cs633
```

# Process Placement

```
pmalakar@csews2:~/class/2020-21-II$ mpirun -np 4 -hosts csews1,csews2 ./3.multipletaskscoreID | sort -k1n
0 of 4 Running on csews1:2
1 of 4 Running on csews2:6
2 of 4 Running on csews1:10
3 of 4 Running on csews2:10
pmalakar@csews2:~/class/2020-21-II$ mpirun -np 4 -hosts csews1,csews2 ./3.multipletaskscoreID | sort -k1n
0 of 4 Running on csews1:2
1 of 4 Running on csews2:3
2 of 4 Running on csews1:7
3 of 4 Running on csews2:5
pmalakar@csews2:~/class/2020-21-II$ mpirun -np 4 -hosts csews1:2,csews2:2 ./3.multipletaskscoreID | sort -k1n
0 of 4 Running on csews1:8
1 of 4 Running on csews1:10
2 of 4 Running on csews2:2
3 of 4 Running on csews2:5
pmalakar@csews2:~/class/2020-21-II$ mpirun -np 4 -hosts csews1:2,csews2:2 ./3.multipletaskscoreID | sort -k1n
0 of 4 Running on csews1:4
1 of 4 Running on csews1:2
2 of 4 Running on csews2:3
3 of 4 Running on csews2:6
```

# Executing MPI programs

- Check `which mpicc` and `which mpiexec`
  - Update PATH environment variable
- Compile
  - mpicc –o filename filename.c
- Run
  - mpiexec –np 4 –f hostfile filename [Often “No such file” error]
  - mpiexec –np 4 –f hostfile ./filename

# Resources for MPI

- Marc Snir, Steve W. Otto, Steven Huss-Lederman, David W. Walker and Jack Dongarra,  
MPI - The Complete Reference, Second Edition, Volume 1, The MPI Core.
- William Gropp, Ewing Lusk, Anthony Skjellum,  
Using MPI: portable parallel programming with the message-passing interface, 3rd Ed., Cambridge MIT Press, 2014.
- <https://www mpi-forum.org/docs/mpi-4.0/mpi40-report.pdf>


# Message Passing Paradigm

- Message sends and receives
- Explicit communication


## Communication patterns

- Point-to-point
- Collective

# Communication – Message Passing



# Message Passing



# MPI Basics

- Process is identified by rank
- Communicator specifies communication context

Message Envelope

Source

Destination

Communicator


Tag (0:MPI\_TAG\_UB)

# MPI Data Types

- MPI\_BYTE
- MPI\_CHAR
- MPI\_INT
- MPI\_FLOAT
- MPI\_DOUBLE

# Point-to-point Communication

**MPI\_Send**




SENDER

Blocking send and receive

```
int MPI_Send (const void *buf, int count,  
MPI_Datatype datatype, int dest, int tag,  
MPI_Comm comm)
```

**MPI\_Recv**



RECEIVER

Tags should match

```
int MPI_Recv (void *buf, int count,  
MPI_Datatype datatype, int source, int tag,  
MPI_Comm comm, MPI_Status *status)
```

# **MPI\_Send** Parameters

## **buf**

initial address of send buffer (choice)



## **count**

number of elements in send buffer (non-negative integer)

## **datatype**

datatype of each send buffer element (handle)

## **dest**

rank of destination (integer)

## **tag**

message tag (integer)

## **comm**

communicator (handle)

[https://www.mpich.org/static/docs/latest/www3/MPI\\_Send.html](https://www.mpich.org/static/docs/latest/www3/MPI_Send.html)

# Example

```
int MPI_Send (const void *buf, int count, MPI_Datatype  
datatype, int dest, int tag, MPI_Comm comm)
```

Send 1 INT from rank 0 to rank 1

```
// Initialization  
  
if (myrank == 0)  
    MPI_Send (buf, 1, MPI_INT, 1, 1, MPI_COMM_WORLD);  
else if (myrank == 1)  
    MPI_Recv (buf, 1, MPI_INT, 0, 1, MPI_COMM_WORLD, &status);
```

# Code

```
#include <stdio.h>
#include <stdlib.h>
#include "mpi.h"

int main( int argc, char *argv[])
{
    int myrank, size;
    MPI_Status status;
    double sTime, eTime, time;

    MPI_Init(&argc, &argv);

    int count = atoi (argv[1]);
    int buf[count];

    MPI_Comm_rank(MPI_COMM_WORLD, &myrank) ;
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    // initialize data
    for (int i=0; i<count; i++)
        buf[i] = myrank+i;

    if (myrank == 0)
        MPI_Send (buf, count, MPI_INT, 1, 1, MPI_COMM_WORLD);
    else
        if (myrank == 1)
            MPI_Recv (buf, count, MPI_INT, 0, 1, MPI_COMM_WORLD, &status);

    printf ("%d %d\n", myrank, count);

    MPI_Finalize();
    return 0;
}
```

# Simple Send/Recv Code (sendmessage.c)

```
//From https://www.mpi-forum.org/docs/mpi-3.1/mpi31-report.pdf Chap-3


#include <stdio.h>
#include <string.h>
#include "mpi.h"

int main( int argc, char *argv[] )
{
    char message[20], recvmessage[100];
    int myrank;
    MPI_Status status;
    MPI_Init( &argc, &argv );
    MPI_Comm_rank( MPI_COMM_WORLD, &myrank );
    if (myrank == 0) /* code for process 0 */
    {
        strcpy(message,"Hello, there");
        MPI_Send(message, strlen(message)+1, MPI_CHAR, 1, 99, MPI_COMM_WORLD);
    }
    else if (myrank == 1) /* code for process 1 */
    {
        MPI_Recv(recvmessage, 15, MPI_CHAR, 0, 99, MPI_COMM_WORLD, &status);
        printf("received : %s\n", message);
    }
    MPI_Finalize();
    return 0;
}
```

No runtime  
or compile-  
time error

```
#include <stdio.h>
#include <string.h>
#include "mpi.h"

int main( int argc, char *argv[] )
{
    char message[20], recvmessage[100];
    int myrank;
    MPI_Status status;
    MPI_Init( &argc, &argv );
    MPI_Comm_rank( MPI_COMM_WORLD, &myrank );
    if (myrank = 0) /* code for process 0 */
    {
        strcpy(message,"Hello, there");
        MPI_Send(message, strlen(message)+1, MPI_CHAR, 1, 99, MPI_COMM_WORLD);
    }
    else if (myrank = 1) /* code for process 1 */
    {
        MPI_Recv(recvmessage, 10, MPI_CHAR, 0, 99, MPI_COMM_WORLD, &status);
        printf("received : %s\n", recvmessage);
    }
    MPI_Finalize();
    return 0;
}
```



Runtime  
error

# Message Size

Sender

message (13 bytes)

Receiver

Message (10 bytes)

```
Fatal error in MPI_Recv: Message truncated, error stack:  
MPI_Recv(200).....: MPI_Recv(buf=0x7ffccc37c610,  
count=10, MPI_CHAR, src=0, tag=99, MPI_COMM_WORLD,  
status=0x7ffccc37c5d0) failed  
MPIDI_CH3_PktHandler_EagerShortSend(363): Message from rank 0  
and tag 99 truncated; 13 bytes received but buffer size is 10
```